

# ICE2607 Lab 4: LSH

SJTU-SEIEE    cny123222

December 19, 2024

## 1 Experiment Overview (实验概览)

**Locality Sensitive Hashing (LSH)** is a technique that facilitates efficient Approximate Nearest Neighbor Search, offering a faster alternative to traditional methods like Nearest Neighbor (NN) or K-Nearest Neighbor (KNN) search. **By grouping data into buckets based on hash functions, LSH significantly accelerates search processes by narrowing down the search space.**

This experiment aims to compare the efficiency of LSH with NN search methods. Using a dataset of 40 images, the goal is to find the most similar image to a target image and analyze the time disparity between LSH and NN searches. The influence of different projection sets on initialization and search speed is also investigated. Furthermore, the experiment explores the use of ResNet for extracting image features to enhance retrieval performance.

## 2 Solution Approach (解决思路)

### 2.1 Step 1: Feature Vector Generation

Firstly, image features are extracted using color histograms. Each image is evenly divided into four parts, and the RGB proportions of these quadrants are concatenated to create a 12-dimensional vector. Through an encoding step, a feature vector consisting of 12 components with values of 0, 1, or 2 is generated for each image. The process, including the encoding rule, is illustrated in the following pseudocode:

```
1 def generate_feature_vector(img): # img.shape = (H, W, 3)
2     H, W, C = img.shape
3     half_H, half_W = H // 2, W // 2
4     # Calculate RGB proportions
5     rgbs = []
6     for i in range(2):
7         for j in range(2):
8             subimage = img[i * half_H: (i + 1) * half_H, j * half_W: (j + 1) * half_W] #
9                 subimage.shape = (H//2, W//2, 3)
10             rgb = np.sum(subimage, axis=(0, 1)) # rgb.shape = (3,)
11             rgbs.extend(rgb / np.sum(rgb))
12     rgbs = np.array(rgbs) # rgbs.shape = (12,)
```

```

12  # Encoding
13  low_th = min(rgbs) + (max(rgbs) - min(rgbs)) / 3 # Lower threshold
14  high_th = max(rgbs) - (max(rgbs) - min(rgbs)) / 3 # Higher threshold
15  feature_vec = np.ones(12, dtype=np.uint8)
16  feature_vec[rgbs >= high_th] = 2 # Assign 2
17  feature_vec[rgbs <= low_th] = 0 # Assign 0
18  return feature_vec

```

## 2.2 Step 2: Hash Function Computation

Next, the 12-dimensional feature vector is mapped to a Hamming space of  $d'$  dimensions and then projected onto a projection set, dividing the  $N$  images into  $n$  groups, where  $n \ll N$ . Since the details of the hash function are intricate, its descriptions are simplified here. However, the implementation, as shown in the following pseudocode, can be concise.

```

1  def hash_proj(feature_vec, proj_set):
2      aux1 = (proj_set - 1) // 2
3      aux2 = (proj_set - 1) % 2 + 1
4      hash_val = (aux2 <= feature_vec[aux1]).astype(np.uint8)
5      return tuple(hash_val) # hash_val will be used as the key in a dict

```

## 2.3 Step 3: LSH Retrieval

After the first two steps, the dataset images are classified into  $n$  categories. When given the target image, its hash value is computed and the image is classified. Then, we only need to check the dataset images in the same category. To find the most similar image, we compute the minimum distance between the feature vectors. The process can be illustrated using the following pseudocode:

```

1  def LSH_retrieval(hash_dict, target_img, proj_set): # hash_dict maps hash values to a
2      list of image paths
3      # Compute hash value of target image
4      target_vec = generate_feature_vector(img)
5      hash_val = hash_proj(target_vec, proj_set)
6      # Find the nearest dataset image
7      candidates = hash_dict[hash_val] # Assume hash_val in hash_dict.keys()
8      result = find_nearest(target_vec, candidates)
9      return result

```

## 3 Experimental Results (实验结果)

### 3.1 Experimental Environment

The experiments were conducted on macOS Sonoma 14.6.1 with Python 3.13.0. Key libraries included OpenCV-Python 4.10.0.84, Numpy 2.1.3 and Matplotlib 3.9.2.

### 3.2 Retrieval Results

Figure 1 illustrates the retrieval results of the LSH and NN algorithms. Both methods consider **38.jpg** as the most similar image, which is identical to the target image. However, LSH demonstrates a **speedup of approximately 12 times** compared to NN, aligning with our expectations. This improvement is attributed to LSH effectively narrowing down the search space, and thus reducing the search time required.

```
LSH:
Projection Set:[1, 8, 16, 24]
Most similar image: dataset/38.jpg
Time taken: 0.002923727035522461s

NN:
Most similar image: dataset/38.jpg
Time taken: 0.035111188888549805s

Speed up: 12.01x
```

Figure 1 Retrieval Results of LSH and NN with Speed Comparison

## 4 Analysis and Discussion (分析与思考)

### 4.1 Influence of Projection Sets

In the previous experiment, the projection set was manually set as [1, 8, 16, 24]. To enhance the efficiency of the LSH method, we delve into the impact of varying projection sets on retrieval speed. The dimensions of the projection set range from 1 to 24, and we scatter the numbers in the projection sets so that they contain more information. For example, with a dimension of 3, the projection set is set as [1, 12, 24].

#### 4.1.1 Influence on Search Time

Figure 2 illustrates how different projection sets influence the speedup of LSH over NN. From the figure, it is evident that as **the dimension of the projection sets  $m$**  increases, the retrieval speed initially experiences a sharp rise and then stabilizes. This pattern is straightforward to comprehend.

- When  $m = 1$ , the dataset is divided into two categories, halving the number of images for comparison and resulting in a 2x speedup.
- As  $m$  increases, the number of categories grows exponentially, reducing the number of candidate images exponentially, thereby sharply increasing the retrieval speed.
- When  $m = 5$ , given that  $\log_2 40 \approx 5$ , there may be only one image in each category, necessitating comparison with just one image against the target image.

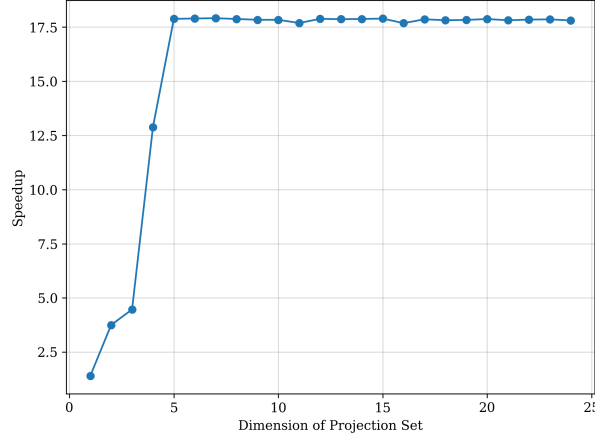


Figure 2 Impact of Projection Sets on Retrieval Speed

- With further increments in  $m$ , where there remains only one image in each category, the situation remains unchanged. LSH reaches the retrieval speed limit at this point.

#### 4.1.2 Influence on Initialization Time

Does a larger value of  $m$  always lead to better retrieval performance? Figure 3 illustrates the impact of  $m$  on both the initialization time and search time of LSH. When search time decreases with increasing  $m$ , as we discussed earlier, the initialization time for LSH, which involves computing hash values and classifying dataset images, increases linearly. The exact reason for this is not clear and could be attributed to the additional time required for computing hash values and inserting them into the hash table.

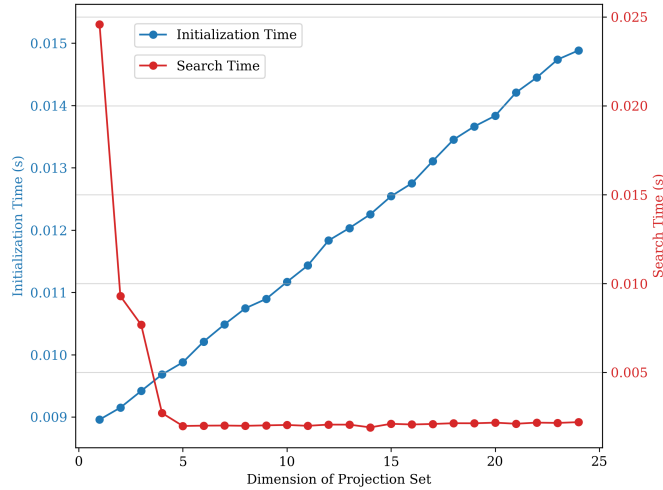


Figure 3 Impact of Projection Sets on Initialization and Search Time

#### 4.1.3 Suggestion for Choosing Projection Sets

From the observations, it is evident that as the dimension of the projection sets increases, the initialization time gradually increases, while the search time first sharply decreases before stabilizing. Therefore, selecting an appropriate projection set is crucial for the efficiency of the LSH algorithm.

Based on our experimental results, we recommend **opting for a uniformly distributed projection set with a dimension of  $\lfloor \log_2 N \rfloor - 1$** , where  $N$  represents the number of images in the dataset. This choice strikes a balance between shorter initialization time and shorter search time. Therefore, in this specific case, we selected the projection set [1, 8, 16, 24].

## 4.2 Feature Analysis and ResNet Feature Extraction

### 4.2.1 Feature Analysis

In the previous experiment, color histograms were employed to extract image features. Figure 4 displays the cosine similarity between image feature vectors, providing valuable insights into the extracted features.

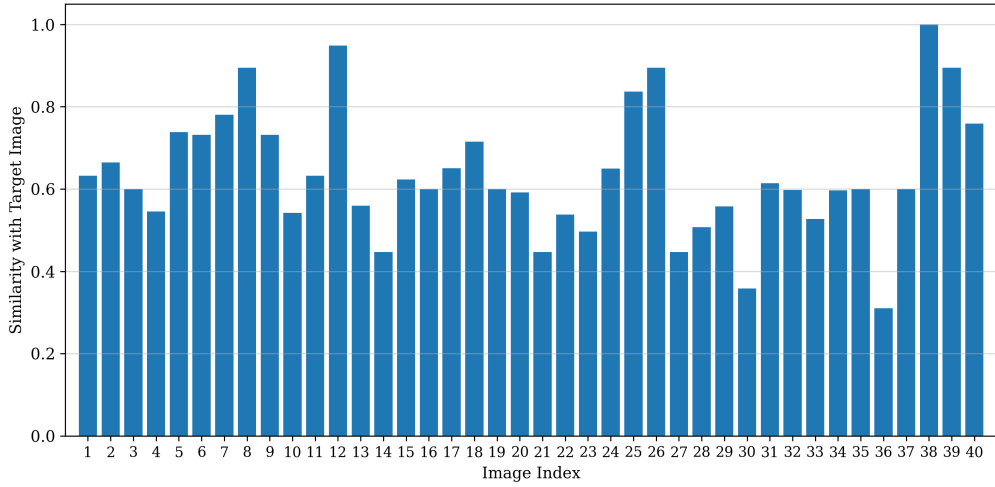


Figure 4 Similarity between Dataset Images and Target Image Using Color Histogram Features

Analyzing Figure 4, it is clear that, besides 38.jpg, images such as 12.jpg, 8.jpg, 26.jpg, 39.jpg, and 25.jpg exhibit a strong similarity with the target image based on the feature vectors derived from color histograms. Referring to Figure 5, it can be noticed that these images share a similar hue of yellow, aligning with our expectations regarding the features extracted by color histograms.

The second most similar image (12.jpg) indeed bears a striking resemblance to the target image, validating the efficacy of the algorithm. Images 26.jpg, 39.jpg, and 25.jpg also present a similar desert background, further confirming the effectiveness of retrieval, particularly concerning similar colors.

However, image 8.jpg, despite having yellow as the predominant color, does not share significant visual similarity with the target image. Consequently, we proceed to explore alternative feature extraction methods.

### 4.2.2 ResNet Feature Extraction

Considering the effectiveness of deep learning models in extracting high-level features, a **pretrained ResNet-18 model** is utilized to extract features from the images. To be consistent with the previous experiment, we employ the first 12 dimensions of the output of the network. Figure 6 illustrates the similarity between the image feature vectors generated by ResNet.

Observing Figure 6, it can be noted that the second most similar image is also 12.jpg, consistent with



38.jpg (Target Image)



12.jpg



8.jpg



26.jpg



39.jpg



25.jpg



2.jpg

Figure 5 Several Images in the Dataset

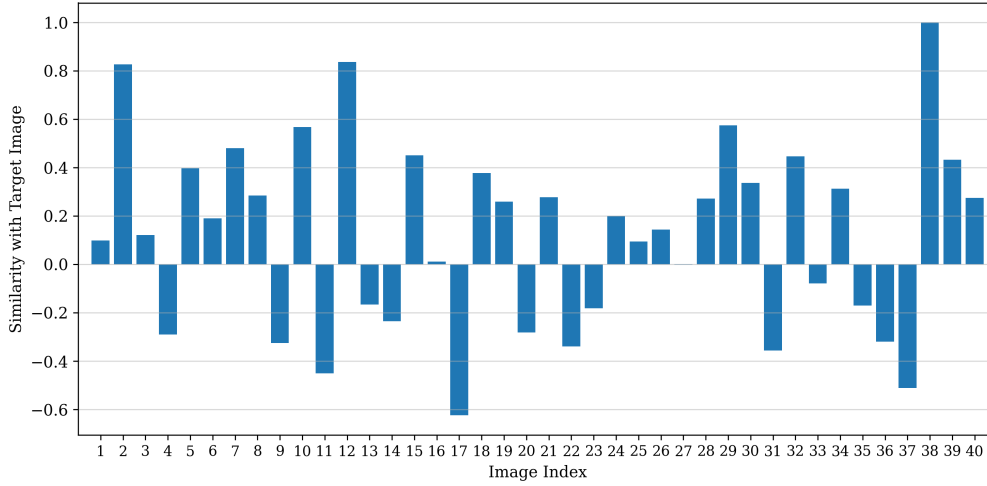


Figure 6 Similarity between Dataset Images and Target Image Using ResNet Features

the results obtained using color histograms. However, the third most similar image is 2.jpg, which features a plant similar to the one in the target image. This highlights the effectiveness of ResNet in extracting high-level features, surpassing the color histograms that primarily focus on color distribution. Furthermore, the feature vectors provided by ResNet exhibit, on average, lower similarity levels than those given by color histograms, which aids in distinguishing between different images.

### 4.3 Highlights of Codes

- **Multithreaded Image Loading:** Employing multithreading for image loading to expedite LSH initialization, mitigating potential I/O bottlenecks.
- **Integration of tqdm Progress Bar:** Incorporating a tqdm progress bar to display real-time progress updates, improving user experience and task monitoring.
- **Object-Oriented Approach:** Utilizing an object-oriented design to enhance code readability and maintainability.
- **Command-Line Interface:** Implementation of a command-line interface allowing for the input of various parameters, enhancing versatility, usability, and cross-system compatibility.

## 5 Reflection and Conclusion (实验感想)

### 5.1 Learning from Experiments

Through the course of these experiments, we have deepened our comprehension of LSH and image retrieval, refined our skills in parameter selection, and further honed our proficiency in LaTeX. These experiences have significantly boosted our research capabilities, proving to be invaluable for those embarking on scientific exploration.

## 5.2 Challenges Encountered

Two main challenges were encountered during the experiments.

- Firstly, the encoding rules provided by the slides mainly resulted in binary encoding (0 and 1) without the possibility of obtaining a value of 2. This limitation led to poor encoding effectiveness, necessitating the exploration of alternative encoding methods for experimentation.
- Secondly, measuring the algorithm's runtime posed difficulties as it required multiple measurements to calculate averages. Each measurement was susceptible to interference from other processes running on the computer, causing abrupt fluctuations in the results. Additionally, this process was time-consuming.

## 5.3 Conclusion

In this experiment, we employed the LSH algorithm to retrieve the target image from the dataset. We investigated the impact of projection sets on retrieval speed and proposed guidelines for selecting suitable projection sets. Furthermore, we analyzed features extracted by color histograms and enhanced them using ResNet.



## A Source Code File List

Table 1 File List

File Name	Description
knn.py	Nearest Neighbor algorithm
lsh.py	Locality Sensitive Hashing algorithm
main.py	Image retrieval process and main function
plot.py	Functions for generating analytical plots
preprocess.py	Image preprocessing functions

## B Source Code

```
1  """
2  File name: knn.py
3  """
4  import numpy as np
5  from preprocess import Image
6
7
8  class KNN:
9      """
10     K-Nearest Neighbors (K = 1)
11     """
12
13     def __init__(
14         self,
15         resnet: bool = False,
16         normalize: bool = True
17     ):
18         self.image_paths = []
19         self.resnet = resnet
20         self.normalize = normalize
21
22     def add(
23         self,
24         img_path: str
25     ):
26         """
27         Add a dataset image
28         """
29         self.image_paths.append(img_path)
30
```

```

31     def search(
32         self,
33         img_path: str
34     ):
35         """
36         Search for the most similar image in the dataset
37         """
38         image = Image(img_path, self.resnet, self.normalize)
39
40         min_dist = float('inf')
41         min_img = None
42
43         for img_path in self.image_paths:
44             img, dist = self._calc_dist(img_path, image)
45             if dist < min_dist:
46                 min_dist = dist
47                 min_img = img
48
49         return min_img.path
50
51     def _calc_dist(
52         self,
53         img_path: str,
54         target_img: Image
55     ) -> float:
56         """
57         Calculate the distance between the target image and the image in the dataset
58         """
59         img = Image(img_path, self.resnet, self.normalize)
60         dist = np.linalg.norm(img.feature_vec - target_img.feature_vec)
61         return img, dist

```

```

1     """
2     File name: lsh.py
3     """
4     import numpy as np
5     from typing import Union
6     from preprocess import Image
7
8
9     class LSH:
10         """
11         Locality Sensitive Hashing
12         """

```

```

13
14     def __init__(
15         self,
16         indicators: Union[list, np.ndarray], # Projection set
17         resnet: bool = False,
18         normalize: bool = True
19     ):
20         self.hash_tab = {} # Hash table to store the dataset
21         self.indicators = np.array(indicators) if isinstance(indicators, list) else
22             indicators
23         self.resnet = resnet
24         self.normalize = normalize
25
26     def add(
27         self,
28         img_path: str
29     ):
30         """
31         Add a dataset image to the hash table
32         """
33         img = Image(img_path, self.resnet, self.normalize)
34         hash_val = self._projection(img.feature_vec, self.indicators)
35         if hash_val not in self.hash_tab:
36             self.hash_tab[hash_val] = [img_path]
37         else:
38             self.hash_tab[hash_val].append(img_path)
39
40     def search(
41         self,
42         img_path: str
43     ):
44         """
45         Search for the most similar image in the dataset
46         """
47         image = Image(img_path, self.resnet, self.normalize)
48         hash_val = self._projection(image.feature_vec, self.indicators)
49         if hash_val not in self.hash_tab:
50             return None
51
52         candidates = self.hash_tab[hash_val]
53         min_dist = float('inf')
54         min_img = None
55
56         for img_path in candidates:

```

```

56         img, dist = self._calc_dist(img_path, image)
57         if dist < min_dist:
58             min_dist = dist
59             min_img = img
60
61         return min_img.path
62
63     def _projection(
64         self,
65         feature_vec: np.ndarray,
66         indicators: np.ndarray
67     ) -> np.ndarray:
68         """
69         Project the feature vector to the given projection set
70         """
71         proj_1 = (indicators - 1) // 2
72         proj_2 = (indicators - 1) % 2 + 1
73         hash_val = (proj_2 <= feature_vec[proj_1]).astype(np.uint8)
74         return tuple(hash_val) # hash_val will be used as the key in a dict
75
76     def _calc_dist(
77         self,
78         img_path: str,
79         target_img: Image
80     ) -> float:
81         """
82         Calculate the distance between the target image and the image in the dataset
83         """
84         img = Image(img_path, self.resnet, self.normalize)
85         dist = np.linalg.norm(img.feature_vec - target_img.feature_vec)
86         return img, dist

```

```

1  """
2  File name: main.py
3  """
4  import sys
5  import glob
6  import time
7  import argparse
8  from lsh import LSH
9  from knn import KNN
10
11
12  def parse_args(args):

```

```

13     """
14     Parse command line arguments
15     """
16     parser = argparse.ArgumentParser()
17     parser.add_argument("--image-dir", type=str, default="dataset", help="Path to folder
18         containing dataset images")
19     parser.add_argument("--type", type=str, choices=["LSH", "NN", "comp"], default="LSH",
20         help="Choose the algorithm to use, comp for comparison")
21     parser.add_argument("--indicator", type=list, default=[1, 8, 16, 24],
22         help="Projection set for LSH")
23     parser.add_argument("--resnet", type=bool, default=False, help="Use ResNet to
24         generate feature vector")
25     parser.add_argument("--target-dir", type=str, default="target.jpg", help="Path to the
26         target image")
27
28     args = parser.parse_args(args)
29     return args
30
31 def get_image_paths(input_dir, extensions = ("jpg", "jpeg", "png", "bmp")):
32     """
33     Get image paths from the given directory
34     """
35     pattern = f"{input_dir}/**/*"
36     img_paths = []
37     for extension in extensions:
38         img_paths.extend(glob.glob(f"{pattern}.{extension}", recursive=True))
39
40     if not img_paths:
41         raise FileNotFoundError(f"No images found in {input_dir}. Supported formats are:
42             {'', '.join(extensions)}")
43
44     return img_paths
45
46 def main(args):
47     args = parse_args(args)
48
49     tasks = ["LSH", "NN"] if args.type == "comp" else [args.type]
50
51     for type in tasks:
52         # Initialize the searcher
53         searcher = LSH(args.indicator, args.resnet) if type == "LSH" else KNN(args.resnet)

```

```

51     # Add images to the searcher
52     img_paths = get_image_paths(args.image_dir)
53     for img_path in img_paths:
54         searcher.add(img_path)
55
56     # Search for the most similar image
57     start_time = time.time()
58     result_path = searcher.search(args.target_dir)
59     finish_time = time.time()
60
61     # Record time taken for each algorithm
62     if args.type == "comp":
63         if type == "LSH":
64             lsh_time = finish_time - start_time
65         else:
66             knn_time = finish_time - start_time
67
68     print()
69     print(f"{type}:")
70     if type == "LSH":
71         print(f"Projection Set:{args.indicator}")
72         print(f"Most similar image: {result_path}")
73         print(f"Time taken: {finish_time - start_time}s")
74
75     # Compare the time taken for LSH and NN
76     if args.type == "comp":
77         speed_up = knn_time / lsh_time
78         print()
79         print(f"Speed up: {speed_up:.2f}x")
80
81
82 if __name__ == '__main__':
83     main(sys.argv[1:])

```

```

1  """
2  File name: plot.py
3  """
4  import sys
5  import time
6  import threading
7  from tqdm import tqdm
8  import numpy as np
9  from typing import Tuple
10 import matplotlib.pyplot as plt

```

```

11 from lsh import LSH
12 from knn import KNN
13 from preprocess import Image
14
15 # Plot settings
16 plt.rcParams['font.family'] = 'serif'
17 plt.rcParams['font.size'] = 11
18 plt.rcParams['axes.labelsize'] = 11
19 plt.rcParams['xtick.labelsize'] = 11
20
21
22 def compute_knn_time(
23     repeat_times: int = 100
24 ):
25     """
26     Compute the average time for KNN search
27     """
28     knn = KNN(normalize=False) # Skip RGB normalization for better result
29     for i in range(1, 41):
30         img_path = f"dataset/{i}.jpg"
31         knn.add(img_path)
32
33     # Compute the time for KNN search
34     knn_time = 0
35     for i in range(repeat_times):
36         start_time = time.time()
37         knn.search("target.jpg")
38         knn_time += (time.time() - start_time)
39     knn_time /= repeat_times
40
41     return knn_time
42
43
44 def compute_lsh_time(
45     indicators: list,
46     repeat_times: Tuple[int, int] = (1, 100),
47     multithread: bool = True
48 ):
49     """
50     Compute the average time for LSH search
51     """
52     # Compute the time for LSH initialization
53     init_time = 0
54     for i in range(repeat_times[0]):

```

```

55     start_time = time.time()
56     lsh = LSH(indicators, normalize=False) # Skip RGB normalization for better result
57     if multithread: # Use multithreading for better performance
58         threads = []
59         for i in range(1, 41):
60             img_path = f"dataset/{i}.jpg"
61             thread = threading.Thread(target=lsh.add, args=(img_path,))
62             threads.append(thread)
63             thread.start()
64         for thread in threads:
65             thread.join()
66     else:
67         for i in range(1, 41):
68             img_path = f"dataset/{i}.jpg"
69             lsh.add(img_path)
70     init_time += (time.time() - start_time)
71     init_time /= repeat_times[0]
72
73     # Compute the time for LSH search
74     search_time = 0
75     for i in range(repeat_times[1]):
76         start_time = time.time()
77         lsh.search("target.jpg")
78         search_time += (time.time() - start_time)
79     search_time /= repeat_times[1]
80
81     return init_time, search_time
82
83
84 def generate_scatter_indicator(
85     bucket_num: int
86 ):
87     """
88     Generate scatter indicators for LSH
89     """
90     assert 1 <= bucket_num <= 24
91     indicators = []
92     indicators = np.linspace(1, 24, bucket_num, dtype=int)
93     return indicators
94
95
96 def plot_init_search(
97     bucket_nums: range = range(1, 25),
98     repeat_times: Tuple[int, int] = (500, 100),

```



```

99     save_path: str = "results/time_init_search.png"
100 ):
101     """
102     Compare the initialization time and search time of different indicators
103     """
104     # Compute the time for LSH initialization and search
105     lsh_init_times = []
106     lsh_search_times = []
107     for bucket_num in tqdm(bucket_nums, desc="Different Dimension of Projection Set"):
108         indicators = generate_scatter_indicator(bucket_num)
109         init_time, search_time = compute_lsh_time(indicators, repeat_times)
110         lsh_init_times.append(init_time)
111         lsh_search_times.append(search_time)
112     lsh_init_times = np.array(lsh_init_times)
113     lsh_search_times = np.array(lsh_search_times)
114
115     # Plot initialization time
116     fig, ax1 = plt.subplots(figsize=(8, 6))
117     ax1.plot(bucket_nums, lsh_init_times, marker='o', label="Initialization Time",
118             color='tab:blue')
119     ax1.set_xlabel("Dimension of Projection Set")
120     ax1.set_ylabel("Initialization Time (s)", color='tab:blue')
121     ax1.set_ylim(min(lsh_init_times) * 0.95, max(lsh_init_times) * 1.05)
122     ax1.tick_params(axis='y', labelcolor='tab:blue')
123
124     # Plot search time
125     ax2 = ax1.twinx()
126     ax2.plot(bucket_nums, lsh_search_times, marker='o', label="Search Time",
127             color='tab:red')
128     ax2.set_ylabel("Search Time (s)", color='tab:red')
129     ax2.tick_params(axis='y', labelcolor='tab:red')
130
131     # Plot settings
132     plt.grid(True, alpha=0.5)
133     ax1.legend(loc='upper left', bbox_to_anchor=(0.1, 0.97))
134     ax2.legend(loc='upper left', bbox_to_anchor=(0.1, 0.9))
135     plt.tight_layout()
136     plt.savefig(save_path, dpi=300)
137     plt.show()
138
139 def plot_lsh_knn(
140     bucket_nums: range = range(1, 25),
141     repeat_times: int = 100,

```

```

141     save_path: str = "results/time_comp.png"
142 ):
143     """
144     Compare the searching speedup of LSH over KNN
145     """
146     # Compute the time for KNN search
147     knn_time = compute_knn_time(repeat_times)
148
149     # Compute the time for LSH search
150     lsh_times = []
151     for bucket_num in tqdm(bucket_nums, desc="Different Dimension of Projection Set"):
152         indicators = generate_scatter_indicator(bucket_num)
153         _, lsh_time = compute_lsh_time(indicators, (1, repeat_times))
154         lsh_times.append(lsh_time)
155     lsh_times = np.array(lsh_times)
156
157     # Compute the speedup
158     speedup = knn_time / lsh_times
159
160     # Plot the speedup
161     fig = plt.figure(figsize=(8, 6))
162     plt.plot(bucket_nums, speedup, marker='o')
163     plt.xlabel("Dimension of Projection Set")
164     plt.ylabel("Speedup")
165     plt.grid(alpha=0.5)
166     plt.savefig(save_path, dpi=300)
167     plt.show()
168
169
170 def plot_similarity(
171     resnet: bool = False
172 ):
173     """
174     Plot the similarity between the target image and all dataset images
175     """
176     # Compute feature vector for the target image
177     target_img = Image("target.jpg", resnet)
178     target_vec = target_img.feature_vec
179     target_vec = target_vec / np.linalg.norm(target_vec)
180
181     # Compute similarity between the target image and all dataset images
182     similarities = []
183     for i in range(1, 41):
184         img = Image(f"dataset/{i}.jpg", resnet)

```

```

185     img_vec = img.feature_vec
186     img_vec = img_vec / np.linalg.norm(img_vec)
187     similarities.append(np.dot(target_vec, img_vec))
188
189     # Plot the similarity
190     fig = plt.figure(figsize=(10, 5))
191     plt.bar(range(1, 41), similarities)
192     plt.xlabel("Image Index")
193     plt.ylabel("Similarity with Target Image")
194     plt.xlim(0, 41)
195     plt.xticks(range(1, 41), fontsize=10)
196     plt.grid(axis='y', alpha=0.5)
197     plt.tight_layout()
198     save_path = "results/similarity_new.png" if not resnet else
199                 "results/similarity_resnet_new.png"
200     plt.savefig(save_path, dpi=300)
201     plt.show()
202
203 def main(args):
204     if args[0] == "time_comp":
205         plot_lsh_knn(range(1, 25), 100, "results/time_comp_new.png")
206     elif args[0] == "time_init_search":
207         plot_init_search(range(1, 25), (500, 100), "results/time_init_search_new.png")
208     elif args[0] == "similarity":
209         plot_similarity(resnet=False)
210     elif args[0] == "similarity_resnet":
211         plot_similarity(resnet=True)
212
213
214 if __name__ == '__main__':
215     main(sys.argv[1:])

```

```

1  """
2  File name: preprocess.py
3  """
4  import cv2
5  import numpy as np
6  import torch
7  import torch.nn as nn
8  import torchvision.transforms as transforms
9  from torchvision import models
10 from PIL import Image as PILImage
11

```

```

12
13 class Image:
14     """
15     Image class
16     """
17
18     def __init__(
19         self,
20         img_path: str,
21         resnet: bool = False,
22         normalize: bool = True
23     ):
24         self.path = img_path
25         img = cv2.imread(img_path)
26         if not resnet:
27             self.feature_vec = extract_color_feature(img, normalize)
28         else:
29             self.feature_vec = extract_resnet_feature(img)
30
31
32 def extract_color_feature(
33     img: np.ndarray,
34     normalize: bool = True
35 ) -> np.ndarray:
36     """
37     Generate color feature vector for the given image
38     """
39     H, W, C = img.shape
40     half_H, half_W = H // 2, W // 2
41
42     # Calculate the sum of RGB values in each quadrant
43     rgb = []
44     for i in range(2):
45         for j in range(2):
46             quadrant = img[i * half_H: (i + 1) * half_H, j * half_W: (j + 1) * half_W]
47             quadrant_sum = np.sum(quadrant, axis=(0, 1))
48             if normalize:
49                 rgb.extend(quadrant_sum / np.sum(quadrant_sum))
50             else:
51                 rgb.extend(quadrant_sum)
52     rgb = np.array(rgb)
53
54     # Generate feature vector
55     lb = min(rgb) + (max(rgb) - min(rgb)) / 3

```

```

56     ub = max(rgb) - (max(rgb) - min(rgb)) / 3
57     feature_vec = np.ones(12, dtype=np.uint8)
58     feature_vec[rgb >= ub] = 2
59     feature_vec[rgb <= lb] = 0
60     return feature_vec
61
62
63 def extract_resnet_feature(
64     img: np.ndarray
65 ) -> np.ndarray:
66     """
67     Generate feature vector for the given image using ResNet
68     """
69     model = models.resnet18(weights=models.ResNet18_Weights.DEFAULT)
70     model.eval()
71     img = preprocess_resnet(img)
72     with torch.no_grad():
73         features = model(img)
74     features = features.squeeze() # Remove the batch dimension
75     feature_vector = features[:12] # Extract the first 12 features
76     return feature_vector.numpy()
77
78
79 def preprocess_resnet(
80     img: np.ndarray
81 ):
82     """
83     Preprocess the image for ResNet
84     """
85     transform = transforms.Compose([
86         transforms.Resize((224, 224)), # Resize to 224x224
87         transforms.ToTensor(),         # Convert to tensor
88         transforms.Normalize(          # Normalize
89             mean=[0.485, 0.456, 0.406],
90             std=[0.229, 0.224, 0.225]
91         )
92     ])
93     img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
94     img = PILImage.fromarray(img) # Convert to PIL image
95     img = transform(img).unsqueeze(0) # Add batch dimension
96     return img

```